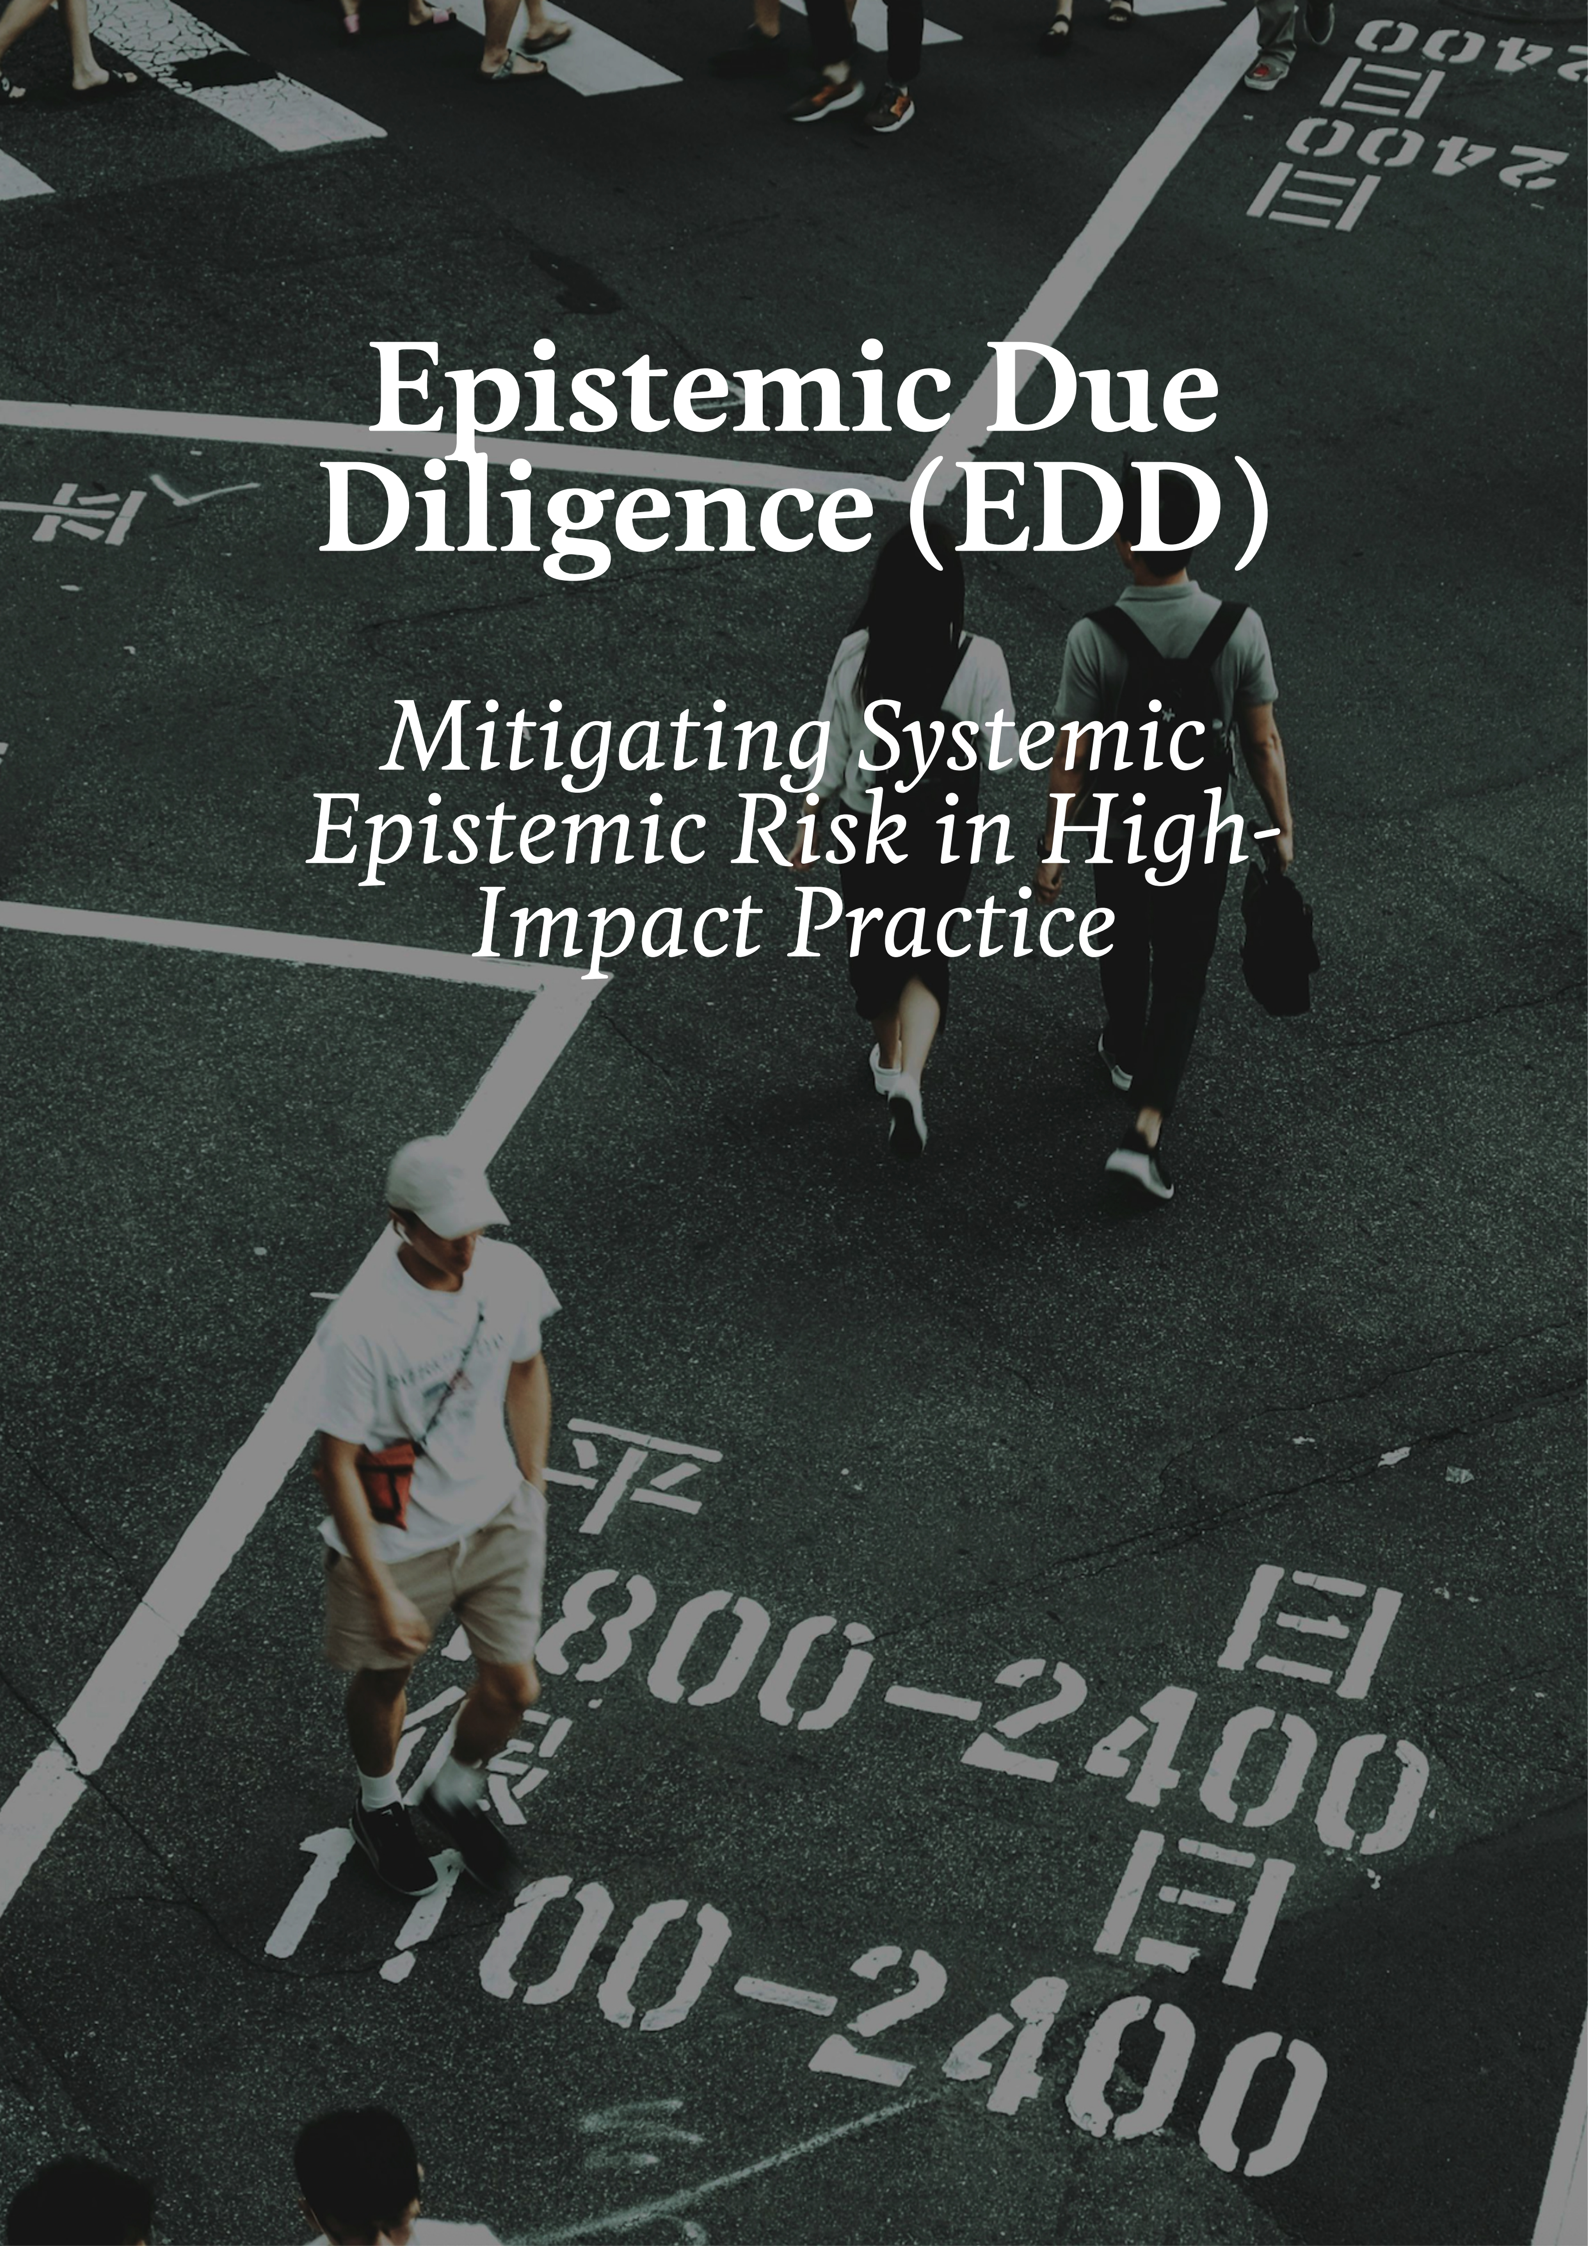

Epistemic Due Diligence (EDD)

Mitigating Systemic Epistemic Risk in High-Impact Practice

Snapshot #1

Imagine a future AI trained on the entire written corpus of the internet—Reddit threads, Wikipedia entries, academic papers, news articles.

Largue Language Models (LLM's) become fluent, omniscient—but also potentially profoundly narrow.

Most of that corpus is in English. Most comes from Western, educated, industrialised contexts. Entire cultures with oral traditions, with ways of knowing that leave no written traces, leave no residue in the training data.

Not because their knowledge is less true. But because it was never written down in a form the crawlers could find.

Here's the problem:

When this system answers questions about "justice" or "the good life," it draws only on what survived.

The rest are absent—not refuted, not considered, but simply never present to be considered.

By the time these systems are deployed globally, they carry a narrowed set of human values. Not because anyone chose that. Because the epistemic filters *(writing, digitisation, English dominance)* did their work upstream, before training began.

What if we asked...

If oral traditions, embodied knowledge, or non-Western conceptual frameworks were excluded, how would we know?

Who decided what counted as "training data"?

What forms of knowledge were structurally incapable of surviving

Could the system ever discover that something is missing?

Or does it only optimise within what it was given?

Imagine...

You have $10 million to improve lives in low-income countries. What do you do? Deworming or cash transfers?

For years, the answer seemed clear. Disability-adjusted life years (DALYs) gave a common currency: "maximise DALYs averted per dollar." Deworming looked extraordinarily cost-effective. By school-attendance records, children were measurably more present.

Cash transfers, harder to measure.

Millions flowed accordingly.

What changed?

The metric.

DALYs measure death and sickness. They do not measure mental health, dignity, or whether people feel their lives are going well. When the *Happier Lives Institute* proposed an alternative *(WELLBYs, which measure life satisfaction directly)* the rankings shifted.

Psychotherapy and cash transfers rose. Some traditional interventions fell.

The shift was not marginal. It was structural. And it happened not because new data arrived, but because someone interrogated the metric. The moral assumptions embedded in the measurement tool had been shaping portfolios all along

The Portfolio Consequence

The DALYs-to-WELLBYs shift is not a one-off correction. It reveals a general vulnerability: any metric can become a bottleneck for value. The question is not whether today's metrics are wrong, but whether we have processes to notice when they are.

What These Examples Share

Both examples follow the same pattern:

A dominant framing shapes how a problem is understood and what counts as a solution. Resources flow accordingly.

But the framing itself rests on assumptions that could be wrong.

And those assumptions were not systematically interrogated before they hardened into practice. Not because anyone was careless. Because the systems we use to evaluate interventions (cost-effectiveness analysis, evidence hierarchies, monitoring frameworks) operate within the chosen frame.

They can tell you how well you are doing against your metrics. They cannot tell you whether your metrics are the right ones.

This is not an edge case. It is a recurring failure mode across global health, development, AI governance, and institutional design. And it has measurable costs.

THINK ABOUT METRICS LIKE:

- **Opportunity cost** *(what else those millions could have funded),*
- **Misallocation** *(resources flowing to less effective interventions),*
- **Goodhart's law / metric capture** *(when the measure becomes the target),*
- **Counterfactual impact** *(the lives not improved because we measured wrong),*
- **Ex ante vs. ex post** *(catching errors before spend, not after),*
- **Robustness / model failure** *(interventions that work on paper fail in reality),*
- **Preventable failure** *(we could have seen this coming),*
- **Scale blindness** *(missing what matters because it's hard to quantify)*
- **Legitimacy critiques** *(are interventions co-created or imposed?)*

Why This Project Exists

We have become exceptionally good at measuring outcomes. But a familiar pattern persists: well-intentioned, evidence-backed interventions repeatedly produce mistrust, iatrogenic harm, or outcomes that "succeed" while corroding legitimacy.

These failures are usually diagnosed late—as implementation problems. But many originate earlier: at the level of problem definition itself.

The costs are not abstract:

Millions misallocated because DALYs made mental health look less cost-effective than deworming—until WELLBYs showed otherwise

Decades of public health damage from France's window tax, which looked like "success" until someone examined the metric

Patients lost to follow-up not because medicine failed, but because their suffering couldn't fit intake forms

In AI governance, the stakes are even larger. If today's epistemic filters shape training data, tomorrow's AI systems will optimise for a narrow subset of human values. Not because anyone chose that. But because the filters did their work upstream.

This project asks: Before we act, whose knowledge shapes how a problem is defined—and whose is filtered out before it can matter at all?

The question is not whether this happens. It is whether we will notice before the costs compound.

Why This Is Worth Funding

IT IS IMPORTANT (I)

The examples above are not hypothetical. They are actual cases where unexamined epistemic filters have:

- Redirected millions away from effective interventions,
- Caused decades of public health damage
- Lost patients who needed help
- Risked encoding a narrowed set of human values into AI systems

This is a recurring failure mode across global health, development, policy, and AI governance. It also addresses a foundational critique of EA: that its commitment to quantification can become blindness—optimising for what is measurable rather than what matters. This project takes that critique seriously, not as an attack but as a question to be answered from within EA's own commitments.

IT IS NEGLECTED (N)

Vast resources flow into measuring outcomes; almost nothing goes into auditing epistemic foundations before outcomes are measured. Stakeholder consultations happen, but they are advisory, not binding. This gap persists not because it is hard to fix, but because no one has built the infrastructure.

IT IS TRACTABLE (T)

The root causes are well-theorised (Fricker, Scott, Han). What is missing is translation into operational tools. A lightweight 3-5 day audit process is entirely feasible—the core questions are knowable, the methods exist, they just have not been systematised into a repeatable framework.

Why EA?

No other community combines rigorous evidence, willingness to interrogate its own methods, cross-cause-area reach, and funding mechanisms that could eventually require EDD as standard practice. This project asks EA to fund the translation of philosophical critique into operational infrastructure—the missing layer between "we have evidence" and "we have the right evidence."

EDD does not reject EA's tools. It extends them—adding an ex ante layer so rigor operates on the fullest possible picture of reality.

For the EA-Community

This is a missing layer
Effective Altruism has given us powerful tools:

Cost-effectiveness analysis

Randomised controlled trials

Evidence hierarchies

Monitoring and evaluation frameworks

These tools are indispensable. But they share a limitation: they operate within an already-chosen frame. They can tell you how well you are doing against your metrics. They can not tell you whether your metrics are the right ones.

Epistemic Due Diligence is not a replacement. It is the missing ex ante layer—an audit of the frame itself before resources are committed. In engineering terms, it is a safety check on your assumptions before you build.

This matters for every cause area:

Global health: Are DALYs capturing what matters, or are WELLBYs revealing something we have missed?

AI safety: Are we solving inner alignment while governance and deployment structures remain under-addressed?

Animal welfare: Whose knowledge counts when designing interventions for factory farms?

Poverty alleviation: What happens when cash transfers look less effective because the metric can not capture dignity?

About Epistemology

Epistemology is simply the study of how we decide what counts as knowledge.

Every institution has built-in assumptions about:

1) what qualifies as evidence,

2) who is considered credible,

3) and which kinds of information are “in scope.”

These assumptions shape how we think about problems before solutions are designed.

When this project talks about “epistemic” failure, it means this:

The system’s way of structuring knowledge filters out important information before action begins.

If key perspectives or risks are excluded at the problem-definition stage, even rigorous, evidence-based interventions can optimise around an incomplete picture of reality.

Have you ever looked at a multiple choice test, and thought:

“None of these are quite right...”?

EXAMPLE

A patient pays regular visits at a clinic with diffuse, cumulative suffering: she describes it as ‘everything just being too much’. She talks about housing insecurity, chronic exhaustion, long Covid, caregiving strain, and economic precarity.

The GP requires discrete, time-bounded symptoms, and has about 10 minutes before his next patient arrives.

Intake forms require discrete, time-bounded symptoms (“sad most days,” “anxious last two weeks”). The patient’s experience does not map cleanly onto available categories.

Clinicians express sympathy but lack procedural tools to register the experience meaningfully. After repeated attempts, the patient disengages and is recorded as “lost to follow-up.

This is not simply misdiagnosis — it is a form of *systemic inadmissibility*.

The system requires suffering to appear in predefined, codifiable categories. We have very little place for chronic or structural stress to go, if it does not fit in a box-to-be-ticked.

When the woman walks away without any help, she feels she has no agency, that this is perhaps her fault, and that care is not able to help her in future instances.

This damage is called iatrogenic harm; harm caused by an intervention that was looking to help.

The Core Diagnosis:
Epistemic Blind Spots Are Structural

Modern institutions rely on what we might call “epistemic filters”, :

1) metrics

2) eligibility criteria

3) evidence hierarchies

4) models and rubrics

These filters are not neutral. They determine what counts as information before any judgment is made.

Philosophical work on epistemic injustice (Miranda Fricker) and legibility (James C. Scott) has shown that:

- some forms of knowledge are routinely excluded not because they are false, but because they are inadmissible

- lived, local, tacit, or ambiguous experience often cannot survive translation into institutional categories

When this happens at scale, the harm is not just mismeasurement—it is epistemic suffering: the erosion of people’s standing as knowers within systems that govern their lives.

Crucially, this exclusion usually happens upstream, before:

- hypotheses are finalised

- metrics are fixed

- or success criteria are locked in

Once that happens, even excellent evaluation tools can only measure within a narrowed frame.

EXAMPLE

In 18th-century France, a tax was levied based on the number of doors and windows in a home—an elegant proxy for wealth. It allowed tax-assessors to measure value without entering homes.

In response, people built houses with fewer windows and doors, or started walling their windows in.

Over time, this led to dark, poorly ventilated homes and long-term public health damage. These damages far outweighted the inital ease of being able to determine taxes from outside of a house.

What went wrong here?

According to James C. Scott's book 'Seeing like a State", this is an example of a system that prioritises the wrong thing: legibility (what can be easily seen), over wellbeing.

The policy:

1) achieved administrative efficiency

2) but álso caused population-level health deterioration

The harm was not a side effect; it was produced by the proxy itself.

Worse, the system could not easily register the harm as policy-induced. The metric looked successful, until the causality was properly examined.

IN LITERATURE

This concept builds on work from multiple traditions. Feminist epistemologies have long documented how women's knowledge must be translated into male-coded language to be taken seriously (Code, 1991, Harding, 1991).

Miranda Fricker's 'Epistemic Injustice' (2007), describes "hermeneutical marginalisation" - when gaps in collective interpretive resources leave marginalised groups unable to make sense of their own experiences in terms institutions understand.

Dorothy Smith's institutional ethnography (1987, 2005) traces how people's everyday lives are "translated" into the abstract categories of bureaucratic texts, a process she calls "textually mediated social organisation."

Empirical research confirms the pattern. In healthcare, studies of patient experience show how clinical intake forms systematically filter out narrative accounts of chronic illness, requiring patients to present discrete, measurable symptoms or be marked as "difficult" or "non-compliant" (Werner & Malterud, 2003; Carel, 2008).

In international development, local knowledge must be rendered into logical frameworks and indicators to access funding; knowledge that cannot be formatted—oral traditions, contextual nuance, embodied practice—is excluded by design (Chambers, 1997; Eyben, 2013).

The burden is not evenly distributed. Those already marginal bear it most heavily. The institution does not adapt; they must. When translation fails—because the categories are too narrow, the cost too high, or the knowledge too tacit to articulate—the institution registers only absence: the patient lost to follow-up, the community that did not apply, the knowledge that left no trace. The failure is invisible because the system cannot see its own filters.

By the time interventions are designed, this work has already happened upstream. The knowledge that remains is the knowledge that fitted the available forms. The rest is absent—not refuted, not considered, but simply too costly to render legible.

Lessons from Critiques of Effective Altruism

These patterns are not abstract. They have played out in movements that aim to do good better.

Effective Altruism (EA) has, in little more than a decade, channelled hundreds of millions toward global health through rigorous cost-effectiveness analysis (MacAskill, 2015; Singer, 2015). The movement has rightly emphasised evidence and scalability. But EA has also attracted sustained critique—critique that illuminates exactly the epistemic exclusion this project addresses.

THE PHILOSOPHICAL CRITIQUE

EA demands absolute impartiality: your uncle's cancer should not matter more than a stranger's malaria, because from "the point of view of the universe" all lives count equally (Singer, 2015; MacAskill, 2015).

Critics argue this abstracts us from the relationships that give moral life meaning (Williams, 1981; Crary, 2021). If moral identities are constituted by particular commitments, then setting them aside may not be neutrality—it may be moral blindness.

African philosophy makes a similar point: "*Ubuntu*" grounds altruism in interdependence, not impartial abstraction (Metz, 2007). Relationality is not an obstacle to doing good but its foundation.

THE INSTITUTIONAL CRITIQUE

The institutional critique. EA's focus on quantifiable, measurable outcomes risks neglecting less-legible work—political advocacy, systemic change, long-term accompaniment (Crary, 2021; Gabriel, 2017). Paul Farmer's work on tuberculosis in Haiti and Peru is instructive: success came not through novel metrics but persistent solidarity with the poor, despite appearing inefficient by standard measures (Farmer, 2001; 2013). A focus on what can be counted may exclude what ultimately counts.

THE EPITEMIC CRITIQUE

EA's knowledge practices—its evidence hierarchies, its metrics, its conception of rigour—privilege Northern, Anglophone, quantitative frameworks (Bhakuni & Abimbola, 2021). In global health, this risks perpetuating the imperial dynamics the field purports to overcome.

Communities in the global South find their knowledge inadmissible, their voices filtered out before problem definition begins (Abimbola & Pai, 2020). The result is not just mismeasurement but epistemic injustice: a wrong done to people as knowers (Fricker, 2007).

The Risk of Double Jeopardy

The stakes are heightened when interventions target vulnerable populations. Global health deals almost categorically with marginalised communities—those already bearing the burden of disease, poverty, and structural violence (Farmer, 2004).

Here, epistemic filters risk creating what has been called *double jeopardy*: the very features that disadvantage a person are subsequently held against them in a different way (Carter, 2001).

A patient with chronic suffering is already disadvantaged by her circumstances. To go back to one of our example, when our patient's experience cannot be translated into discrete clinical categories, she faces a second disadvantage: she is marked as *"lost to follow-up,"* and the system registers no failure. An intervention that succeeds by its own metrics may do so while further entrenching the conditions that produced the problem.

Epistemic Due Diligence will be designed to catch this pattern upstream: before the metrics are locked in, before the patients are lost, before efficiency becomes a justification for exclusion.

From Critique to Construction

This project does not claim to have invented the diagnosis of epistemic exclusion. It draws explicitly on feminist epistemology, institutional ethnography, global health scholarship, and critiques of Effective Altruism.

What it seeks to add is an operational question:

"Can we translate these well-theorised failure modes into a lightweight, repeatable audit that happens before resources are committed and problem definitions are locked in?"

If the critiques are right about what goes wrong, then the responsible next step is to ask what might be done about it—provisionally, testably, and with humility about the limits of any tool.

Epistemic Due Diligence may be that next step.

Why Existing Methods Don't Catch This

This project does not argue that current evidence-based methods are wrong.

Randomised Control Trials (RCTs), cost-effectiveness analysis, Theory of Change, and Monitoring, Evaluation, and Learning (MEL) frameworks are indispensable.

But they share a structural limitation:

They operate ex post, or within an already-defined epistemic frame.

What they do not systematically do is audit:

- how the problem itself was framed

- which forms of knowledge shaped that framing

- which were excluded by design

Stakeholder consultations and participatory methods gesture in this direction, but they are often:

1. *advisory rather than epistemically binding*

2. *unevenly implemented*

3. *not designed to interrogate knowledge hierarchies themselves*

As a result, entire classes of risk remain unmodelled—not because they are rare, but because the system is not built to see them.

BOTTLENECK ANALYSIS

Q: Do we understand the causes of the problem?

No
(Research the problem more)

YES.
The problem is well-diagnosed (Fricker, Scott). The missing piece is operational: no standardised ex ante tool to audit epistemic gaps before resources commit.

Q:
Are there existing solutions with evidence of effeciveness?

Yes

NO.
While we have robust tools to measure outcomes after interventions launch *(RCTs, cost-effectiveness analysis, MEL frameworks)*, there is no standardised ex ante process for auditing whether the problem was framed correctly in the first place. Stakeholder consultations gesture at this but are advisory, uneven, and rarely interrogate knowledge hierarchies themselves.

The bottleneck is not evidence. It's epistemic due diligence before evidence is gathered.

THE BURDEN OF LEGIBILITY

Institutions must simplify reality to function. James C. Scott's 'Seeing like a state' (1998) demonstrates how states impose standard categories (property laws, land registers, censuses) to make populations legible and governable. This simplification is what allows administration, but it also erases complexity. Local knowledges, tacit practices, and informal arrangements disappear from view because they cannot be represented in official formats.

Scott's analysis stops at the state's need for legibility. What remains under-examined is who bears the cost of becoming legible?

I call this the burden of legibility: the work required of individuals and communities to translate their knowledge, experience, or suffering into categories institutions recognise.

EXAMPLE	**THE FILTER**	**THE BURDEN**
Patient lost to follow-up	The intake form requires discrete symptoms the patient is able to define themselves.	Patient bears the cost of translating diffuse suffering into tick-boxes, then walks away unheard
Window tax	Doors/windows as a proxy for wealth	Poor families bear health costs of living in dark, airless homes
DALYs	Health-adjusted life year as a metric	People with mental illness bear the cost of being undervalued in funding allocations
Training data	Written, English-language corpus.	Oral cultures bear the cost of being absent from AI's "values"

The Proposal:
Epistemic Due Diligence (EDD)

Epistemic Due Diligence (EDD) is a proposed ex ante audit designed to fill this gap.

The core idea is straightforward:

Just as we perform due diligence on finances, risks, and outcomes, we should perform due diligence on the epistemic foundations of high-impact interventions—before resources are committed.

EDD is conceived as:

- designed as a bounded, early-stage audit — *modest in time cost, but positioned where changes are cheapest to make.*

- supplementary (not a replacement for quantitative rigor)

- ex ante (before core assumptions harden)

At its simplest, Epistemic Due Diligence asks four questions:

1. KNOWLEDGE INCLUSION

"Who shaped the problem definition?"

2. EPISTEMIC EXCLUSION

"What kinds of knowledge were excluded or made inadmissible?"

3. DISSENT AND NEGATIVE SPACE

"Where does informed dissent exist, and how is it treated?"

4. REVISION PATHWAYS"

"Can excluded knowledge revise the model later—or is it locked out?"

The aim is not consensus, but epistemic robustness.

THEORY OF CHANGE (ToC)

The linear logic from EDD to improved outcomes contains several underspecified mechanisms. Each research question represents not an assumption but an empirical question that this research will investigate. The thesis does not assume these mechanisms work—it asks how, when, and under what conditions they might.

IF... We can make visible the epistemic filters that shape problem definitions before interventions are designed...	**RESEARCH QUESTION 1** How can epistemic exclusion in problem definition be empirically investigated when excluded knowledge leaves no direct trace in institutional records?
THEN... Institutions will more easily recognise when their framing excludes important knowledge...	**RESEARCH QUESTION 1** **Under what conditions do institutions recognise epistemic exlcusion?**
WHICH WILL LEAD TO... Institutions revising problem definitions to be more epistemically inclusive...	**RESEARCH QUESTION 2** **What makes surfaced exclusion lead to revision vs. Dismissal?**
WHICH WILL RESULT IN... Interventions that are more robust, less prone to iatrogenic harm, and defendably more	**RESEARCH QUESTION 3** **How does revised definition translate to design?**
AND ULTIMATELY... Reduced iatrogenic harm, increased institutional legitimacy, new forms of stakeholder inclusion.	**RESEARCH QUESTION 4** **What is the relationship between epistemic inclusion and better outcomes?**

What Is Actually New Here

EDD does not aim to be “more consultation” and not “just an ethics add-on.”

Its innovation lies in three moves:

1. Shifting attention upstream

2. From outcomes and implementation to problem definition itself.

3. Auditing epistemic composition, not opinions

The question is not “what do stakeholders think?” It is:

- “What kinds of knowledge are structurally allowed to count?”

- Translating philosophical critique into operational infrastructure

- Turning well-theorised epistemic failure modes into concrete, repeatable safeguards.

In EA terms, EDD is best understood as a missing layer of ex ante epistemic risk assessment—analogous to safety audits in engineering.

Relevance to AI and Long-Term Governance

This project speaks to AI governance and longtermism in a narrow, disciplined way.

As institutions increasingly rely on algorithmic systems trained on historically filtered data, epistemic exclusion risks becoming infrastructural. Systems optimised for legibility and processability normalise the idea that what cannot be modeled does not matter.

The claim here is modest but has large implications:

If the way we define problems systematically filters out certain perspectives, then any system trained on those definitions will reflect a narrowed set of human values. In AI contexts, this means models may optimise well against what is represented, while missing values that were excluded upstream. EDD does not solve alignment, but it helps reduce the risk that learning starts from an already incomplete picture.

EDD does not solve AI alignment.

It reduces one specific, under-addressed risk: value narrowing before learning even begins.

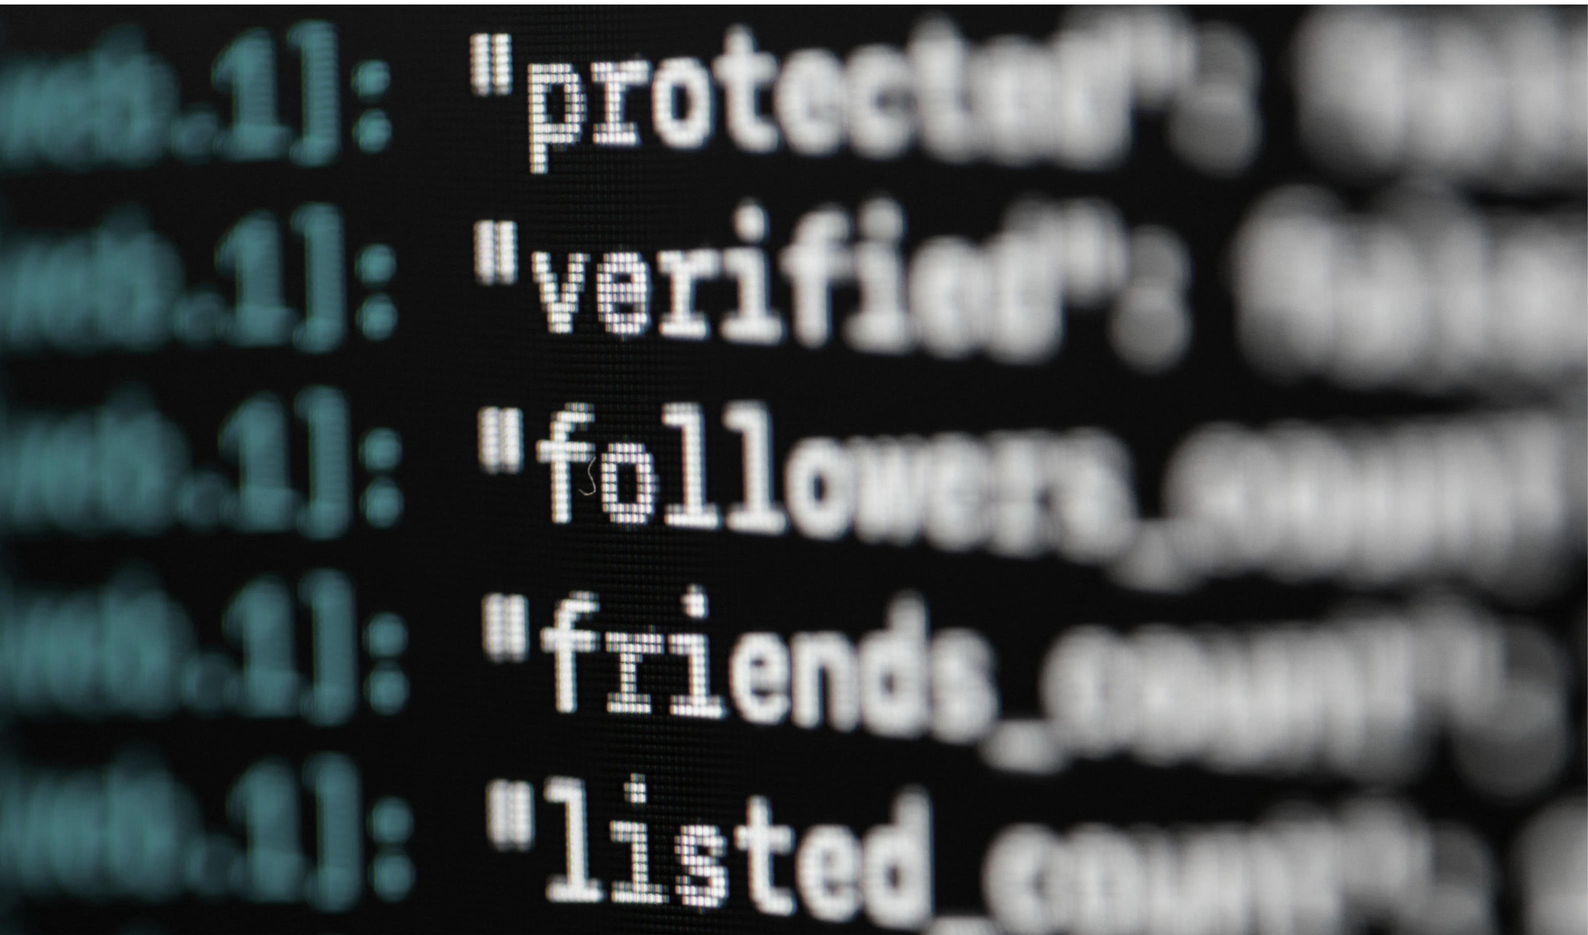

What EDD Could Look Like

Epistemic Due Diligence will be designed as a lightweight, bounded audit—think days, not months. For a hypothetical $10M global health intervention, it might involve:

Step 1: Knowledge Mapping (1-2 days)

Trace who shaped the problem definition. Was it donors? Technical experts? Community representatives? Ministry officials?

List which knowledge forms were admissible (RCT data? Clinical trials?) and which were structurally excluded (oral histories? lived experience? traditional healing practices?)

Step 2: Negative Space Search (1 day)

Identify where informed dissent exists. Are there community members saying "this will not work here"? Are there front-line workers whose daily experience contradicts the model?
Ask: How is this dissent currently treated? As valuable feedback? Or as resistance to be managed?

Step 3: Epistemic Stress-Test (half day)

"If an oral culture with no written language held crucial knowledge about this problem, would our process capture it?"
"If the metric is wrong, would we know? Or would we only optimise within it?"

Step 4: Revision Pathway Check (half day)

Can excluded knowledge revise the model later? Or is the design locked in once funding flows? Where are the feedback loops that don not just measure outcomes, but can challenge the framing itself?

The output is not consensus. It is epistemic transparency—a clear map of what has been included, what has been excluded, and whether the exclusions matter.

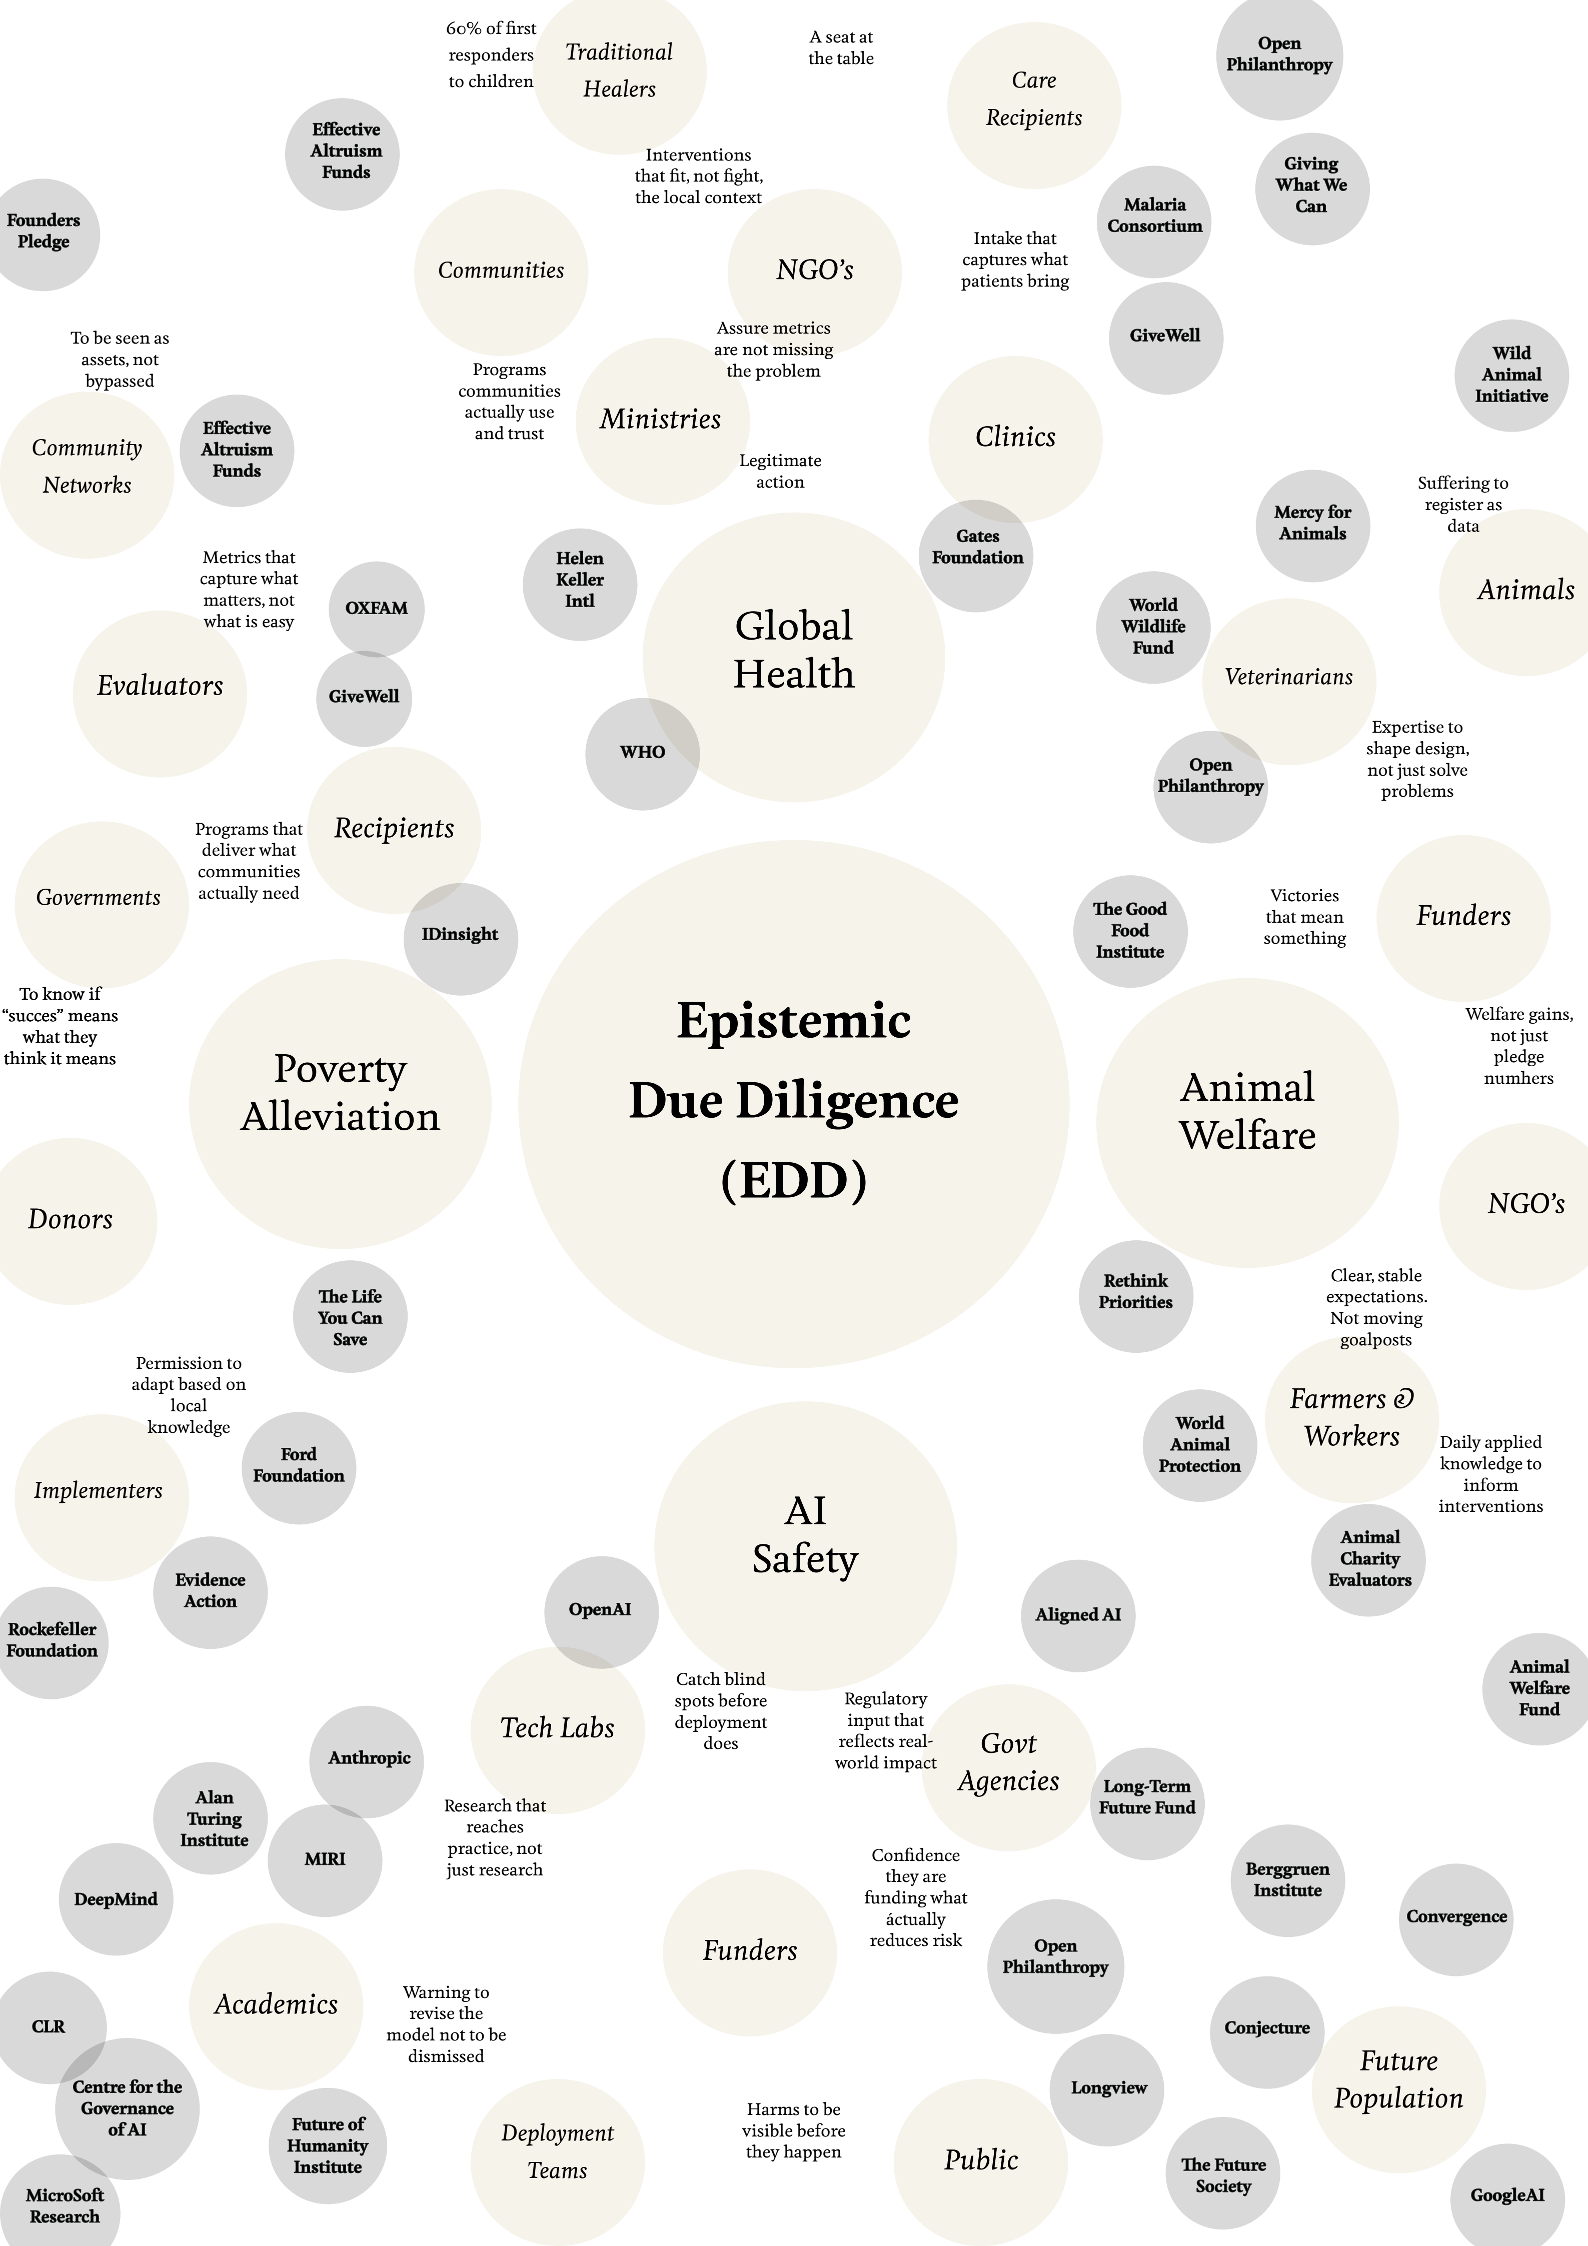
Epistemic Due Diligence (EDD)
Global Health
Traditional Healers
60% of first responders to children
A seat at the table
Care Recipients
Open Philanthropy
Giving What We Can
Malaria Consortium
Effective Altruism Funds
Founders Pledge
Interventions that fit, not fight, the local context
Communities
NGO's
Intake that captures what patients bring
GiveWell
To be seen as assets, not bypassed
Programs communities actually use and trust
Assure metrics are not missing the problem
Ministries
Wild Animal Initiative
Community Networks
Effective Altruism Funds
Legitimate action
Clinics
Suffering to register as data
Mercy for Animals
Gates Foundation
Helen Keller Intl
Metrics that capture what matters, not what is easy
OXFAM
Animals
World Wildlife Fund
Evaluators
GiveWell
Veterinarians
WHO
Open Philanthropy
Expertise to shape design, not just solve problems
Recipients
Programs that deliver what communities actually need
Governments
IDinsight
The Good Food Institute
Victories that mean something
Funders
To know if "succes" means what they think it means
Poverty Alleviation
Animal Welfare
Welfare gains, not just pledge numhers
Donors
NGO's
The Life You Can Save
Rethink Priorities
Clear, stable expectations. Not moving goalposts
Permission to adapt based on local knowledge
Farmers & Workers
World Animal Protection
Daily applied knowledge to inform interventions
Implementers
Ford Foundation
AI Safety
Animal Charity Evaluators
Evidence Action
OpenAI
Aligned AI
Rockefeller Foundation
Animal Welfare Fund
Catch blind spots before deployment does
Regulatory input that reflects real-world impact
Tech Labs
Govt Agencies
Anthropic
Alan Turing Institute
Long-Term Future Fund
Research that reaches practice, not just research
MIRI
Berggruen Institute
Confidence they are funding what áctually reduces risk
DeepMind
Convergence
Funders
Open Philanthropy
Academics
Warning to revise the model not to be dismissed
CLR
Conjecture
Future Population
Centre for the Governance of AI
Longview
Future of Humanity Institute
Deployment Teams
Harms to be visible before they happen
Public
The Future Society
MicroSoft Research
GoogleAI

Known Risks and Open Questions

This is an exploratory project, and several questions remain intentionally open:

1. How minimal can EDD be while still being useful?

2. When does epistemic due diligence add clarity, and when might it add friction?

3. How do we distinguish genuine epistemic blind spots from unavoidable uncertainty?

4. How can this be translated into tools usable by people who are not philosophers?

These are not problems to be solved in advance. They are precisely what this exploratory phase is meant to surface.

What This Stage Is For

The current phase of this project is not about scaling or enforcement.

It is about:

1. clarifying the core insight

2. testing whether the diagnosis resonates across domains

3. stress-testing the proposal with critical, good-faith readers

4. and identifying where this framework is genuinely useful—or not

I am currently sharing this document for reflection, not endorsement.

What Success Could Look Like

If Epistemic Due Diligence becomes standard practice, here is what could change:

For funders:

You don't just ask "What is the cost per DALY?" You also ask "Whose knowledge shaped this DALY calculation? What forms of wellbeing might it miss?"

For implementers:

Before you lock in your theory of change, you audit whose voices were in the room—and whose were structurally incapable of being there.

For the field:

We stop treating "community mistrust" or "implementation failure" as surprises. They become predictable risks we assessed before they happened.

For AI governance:

We do not just ask "How do we align AI with human values?" We first ask "Whose values are currently legible to our systems? Whose are filtered out by the training data itself?"

This is not about slowing down. It is about building interventions that are epistemically robust—grounded in a fuller map of reality, not just the parts that survived our filters. If anything, EDD is about getting it ríght more quickly.

THE STUDY OF ABSENCE

The central methodological challenge of this research is that epistemic exclusion, by definition, leaves no direct trace. Knowledge that was filtered out at the problem-definition stage does not appear in theories of change, funding proposals, or evaluation reports. It is not documented because it was never admitted.

This research therefore does not attempt the impossible task of "finding" excluded knowledge. Instead, it develops and tests methods for studying the operation of filters through indirect evidence. Further research could offer:

1) **A TYPOLOGY OF TRACES** - what kinds of evidence can surface epistemic filters

2) **A METHODOLOGICAL PROTOCOL** - for investigating problem definition retrospectively

3) **EPISTEMIC REFLEXIVITY GUIDES** - for researchers studying their own institutional contexts

BOUNDARY DISCOURSE

"How do institutions justify what they exclude?"

MISFIT EVENTS

"Where do interventions encounter phenomena they were not designed for?"

COUNTERFACTUAL COMPARISON

"How are similar problems framed elsewhere?"

TESTIMONIAL RECONSTRUCTION

"What do those who tried to participate report about barriers?"

About the Researcher

I am an applied philosopher and practice-based researcher working at the intersection of philosophy, institutional life, and framework design. My work addresses a defining condition of contemporary institutions: inadmissibility—the systemic failure to register human complexity within bureaucratic and data-driven systems. I develop practical, philosophically grounded tools that make these pressures legible and open space for agency within them.

This project grows directly from that commitment. Epistemic Due Diligence translates philosophical critique into operational infrastructure—exactly the kind of translation I have spent a decade practicing.

RECENT WORK INCLUDES:

- *Designing integrity notification frameworks for higher education governance*
- *Founder of Applied Olfactory Research Lab with the Royal Academy of Art (KABK)*
- *Designer of ontological meaning-making interventions*
- *Architecting and co-founding the Cambridge Applied Philosophy Society (CAPS)*
- *Academic/corporate retreat facilitator*
- *Advising boards on ethics, integrity, and public values (2019–present)*
- *Co-organising the First Annual Oxbridge Philosophy Row (May 2026)*

ACADEMIC BACKGROUND:

- *University of Cambridge: Advanced Undergraduate Diploma in Research (High First, 81%) with thesis on "Ritual as Technology"*
- *University of Cambridge: Advanced Undergraduate Diploma in Research Theory and Practice (First Honours)*
- *University of Cambridge: Undergraduate Diploma Creativity Theory, History, and Philosophy (First Honours)*
- *Currently: University of Cambridge: Certificate in Philosophy (Pending: First Honours Average)*
- *Effective Thesis Accelerator Fellow (2026)*
- *EA Introductory Course (completed Jan 2026)*

This project does not ask me to become something new. It asks me to do what I already do: translate philosophical critique into tools that work.
(references available upon request)

Exploratory Stakeholder Interview - Michelle Halsall

Founder Beyond the Static & Hearth Leadership
MA Urban Regeneration,
Diploma in Creativity Theory, History & Philosophy

Bali, Indonesia

To start, you shared a powerful comparison between visiting an elder in Bali versus a typical GP visit in the West. Could you paint a picture for us of what that first visit to an elder is actually like? What does a person feel, see, and experience?

Michelle: “On Bali, instead of going to a General Practitioner, people visit an elder, but the role these people have is not easily comparable to the system in the West. They are perhaps best compared to priests, or spiritual advisors or guides. Most villages have one, and you can visit them with anything - every kind of problem.

These elders will first likely explore the problem from a spiritual level. *“What needs to be cleansed, what needs to be purified, what needs to be removed, or taken away from your life.”* A lot of the problems seem to be that people are feeling heavy, or crowded, so a lot of cleansing goes on.

From there, you might be looking more at the emotional side of things: relationships, family, community dynamics. *“What is going on between husband and wife, parents, family, and children?”* Then you might consider the physical. *“What is showing up in your body?”*, and then the most specific problems - the medical.

After all of the clearing and exploring the holistic side of things, what is left, could be a mental or medical issue, that you might want to seek specialist help with. It is a different approach from the West. When you describe the 10 minute GP visits in the Netherlands, I experience a sense of culture shock.”

“When we talk about optimisation, we have to be willing to ask "optimised for what?” ”

The process you describe is fascinating—starting with the spiritual, moving to the emotional, then the physical, and only then considering a medical issue. What is the wisdom in that order? Why start with the "murky" layers?

Michelle: "On Bali, the process always starts from "how you feel". A lot of time is dedicated to exploring what is on the surface, that needs to be removed.

When that first layer of emotional murkiness is resolved, we are able to see what is below that. *"Aah, so below your tiredness may lie the fact that you have been having a lot of arguments with your husband. We might need a family intervention to solve that."*

When thát situation is then resolved, and a problem persists, the next question is: *"So nów what is in here?"* - it is that kind of approach. You don't have to be specific. You just go in and say: *"Things are not quite right?""*

You mentioned that on Bali, the day is structured around ritual—from daily offerings to a day for books, and even a day for machinery. How does that constant, tangible practice of ritual shape a person's sense of purpose and connection to the world?

Michelle: "When it comes to meaning-making on Bali, the culture is influenced by Hinduism, so a lot of the day is structured around ritual. You have daily rituals, weekly rituals, full moons, new moons. Everything is marked - there are occasions to celebrate purpose every day.

Even when people are cooking, a small amount of rice is put on a banana leaf to offer back to nature in gratitude for giving us so much. There is a day for books, every six months, and a day for the motors and machinery that help us do all the work. People here make meaning of all of these little things. You can even tell by the cars which ones have recently been blessed."

In your work at 'Beyond the Static' and 'Hearth Leadership', you talk about restoring energy and trust. How do you view our job of regeneration on a larger scale?

Michelle: "Thinking about regeneration, I would say the job we have to do is twofold: when we focus on just building fairness on a systemic level, we risk losing touch with what makes the day-to-day meaningful. Some days I feel I have been working so hard on creating program, summits, and structures, that I realise I forget I also have to put my feet on the grass. To connect with nature. We need to do both of them.

I think the biggest changes for regeneration come from the human system - from inside of us. When we don't over-intellectualise, some insights come into use when we connect to the world from a heart-level, instead of just from our heads.

For example, when we think about trust on a corporate level, a question we can ask is: "Well, how does trust come into you on a personal level?" That is where we can see those mechanics unfold."

HEARTH
www.hearthleadership.com

MICHELLE HALSALL

DR. CLARE SARAH GOODRIDGE

You describe your partnership with Dr. Clare Sarah Goodridge at Hearth Leadership as a very special connection. Could you tell me more about that?

Michelle; At Hearth Leadership, I work with Dr. Clare Sarah Goodridge, who really excels at working in systems of governance and structure. That is a very fortunate position for someone like me to to be in, because it allows me to focus on the embodied knowledge we gain from listening to our body's perspective. From being creative, and in tune with the world around us. It is almost as if we are two sides of the same brain; the analytical and the intuitive.

You touch on the difficulty of trying to be curious and creative while simultaneously trying to meet rigid system targets. Using the example of a doctor with only 10 minutes, what is lost when we force these two functions into the same tiny box?

Michelle: "When I think of the Western healthcare system, where doctors and nurses generally have very little time, and a lot of output is tied to rules and regulations, that really strikes me.

Doing those two things at once - thinking curiously and creatively, and meeting system targets are very different functions - so they are very hard to switch between and unite in a short span of time. Putting people into systems like that might seem like we are optimising, but if we want to optimise our systems, we have to be willing to ask *"optimised for whát?"."*

You end with a powerful question: When we talk about optimisation, we have to ask "optimised for what?" If we were to optimise a system for wisdom instead of just for speed, what would be the first thing we would need to change?

"*Michelle:* It seems like the Western system that is being exported all over the world is just drifting to a logic of performance - to just one concept of what knowledge is.

I think making space for more of the human system to shine through, and for people to spend more time figuring out what the 'real problem' is, could make us so much more effective at actually solving those problems.

Being curious and taking more than one 'logic' into account ultimately makes us more efficient, expands our capacity and helps us stay grounded in our complex and fast-changing world. To me, that seems more efficient than anything!"

QUESTION 2
Thinking about the last major project you were involved in, I'm going to read five statements. For each one, please rate your agreement on a scale of 1 to 7, where 1 is 'strongly disagree' and 7 is 'strongly agree.' Feel free to explain your rating

	Not at all 1	2	3	4	5	6	*Very much so* 7
"The people who defined the problem represented the full diversity of those affected by it."	o	o	o	o	o	o	o
"There was knowledge about this problem that mattered for the solution but simply could not fit into project's categories or metrics."	o	o	o	o	o	o	o
"Informed dissent *(people who fundamentally challenged our approach)* was actively sought out and taken seriously."	o	o	o	o	o	o	o
"Once our metrics were set, it would have been very difficult to change them based on new learning."	o	o	o	o	o	o	o
"I have concerns that we optimised for what we could measure, but may have missed what ultimately matters."	o	o	o	o	o	o	o

QUESTION 3
Here are five possible outcomes of an Epistemic Due Diligence process. If you could only achieve three of them, which would you choose? And which one is the least important?"

1) **Inclusion:** All relevant perspectives were present at problem-definition.

2) **Revisability:** The project can change its core metrics if early learning suggests it's needed.

3) **Legitimacy:** Affected communities trust that their knowledge was taken seriously.

4) **Efficiency:** The process did not significantly slow down decision-making.

5) **Effectiveness:** The intervention achieved its intended outcomes.

Invitation

If you are reading this, I would value your thoughts on any of the following:

1. Does this identify a real gap you recognise in your own work?

2. Where does the argument feel strongest—or weakest?

3. What existing practices does this overlook or duplicate?

4. Under what conditions would this be genuinely helpful—or genuinely unhelpful?

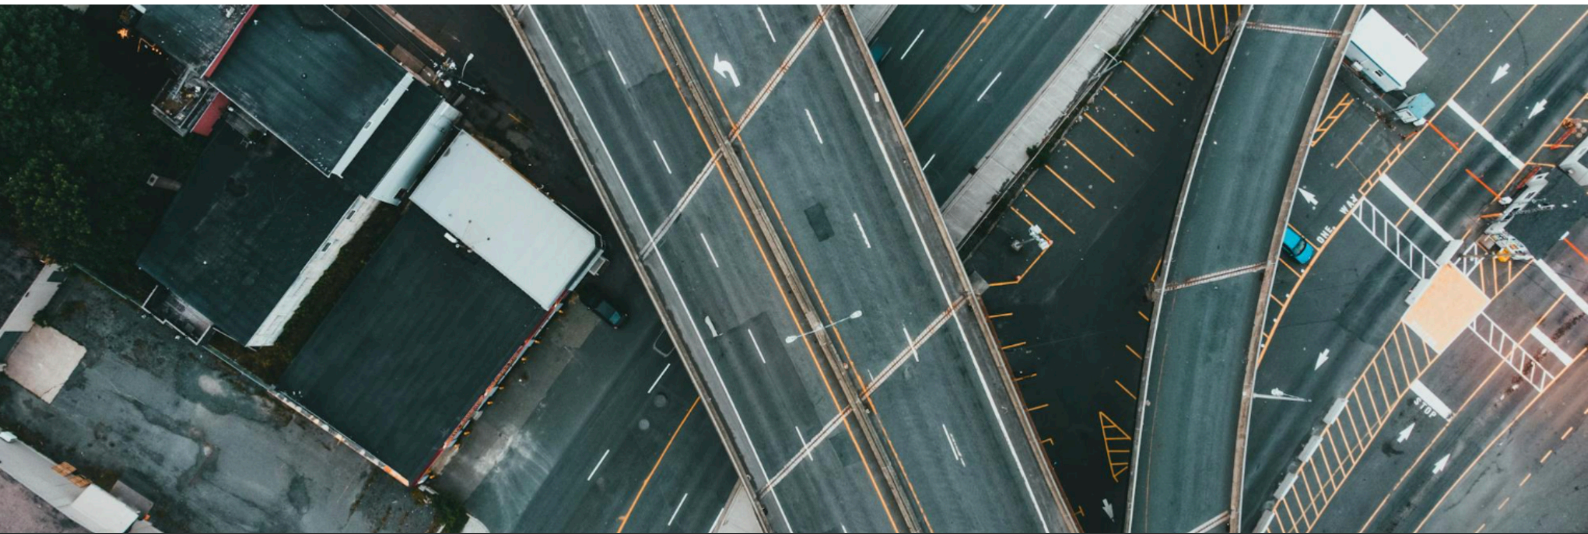

"Epistemic Due Diligence is an attempt to turn a familiar discomfort—"something important is missing here"—into a tractable, testable intervention.

Whether it succeeds depends on critique as much as on agreement.

Thank you for reading!"

Renske van Vroonhoven

rtv22@cam.ac.uk